

COSI 131a: Operating Systems

Assignment Project Exam Help

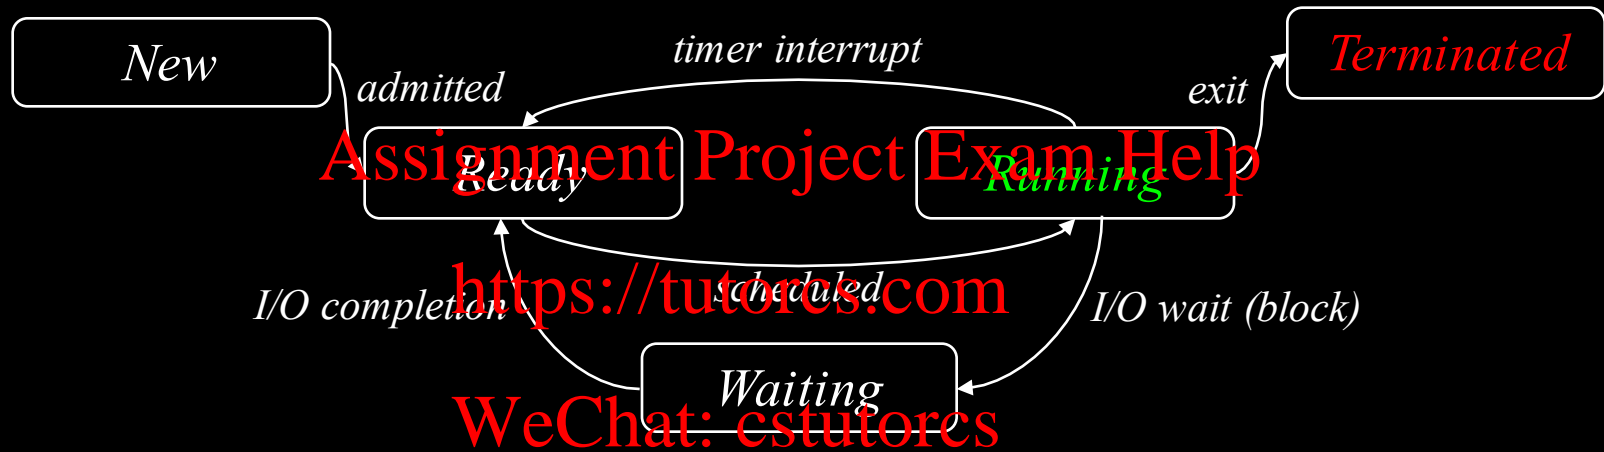
<https://tutorcs.com>

WeChat: cstutorcs (1)

Chapter 6

Review: Process States

Processes Not Always Running



New:

Process is being created

Running:

Instructions are being executed

Waiting:

Process is waiting for some event to occur

Ready:

Process is waiting to be assigned the CPU

Terminated:

Process has finished execution

Agenda

1. Scheduling Overview

- ➡ 1. *First Example of Resource (CPU) Management (Sharing)*
- 2. *Non-Preemptive (N) vs. Preemptive Scheduling (P)*
- 3. *Metrics: Ways to Assess Effectiveness of Scheduling Policies*

Assignment Project Exam Help

2. Scheduling Policies

- 1. *First-Come-First-Served (FCFS) (N)*
- 2. *Shortest-Job-First (SJF) (N, P)*
- 3. *Priority (N, P)*
- 4. *Round-Robin (RR) (P)*
- 5. *Multilevel Queues (MLQ) (P) , Lottery*
- 6. *Real-time*

<https://tutorcs.com>

WeChat: cstutorcs

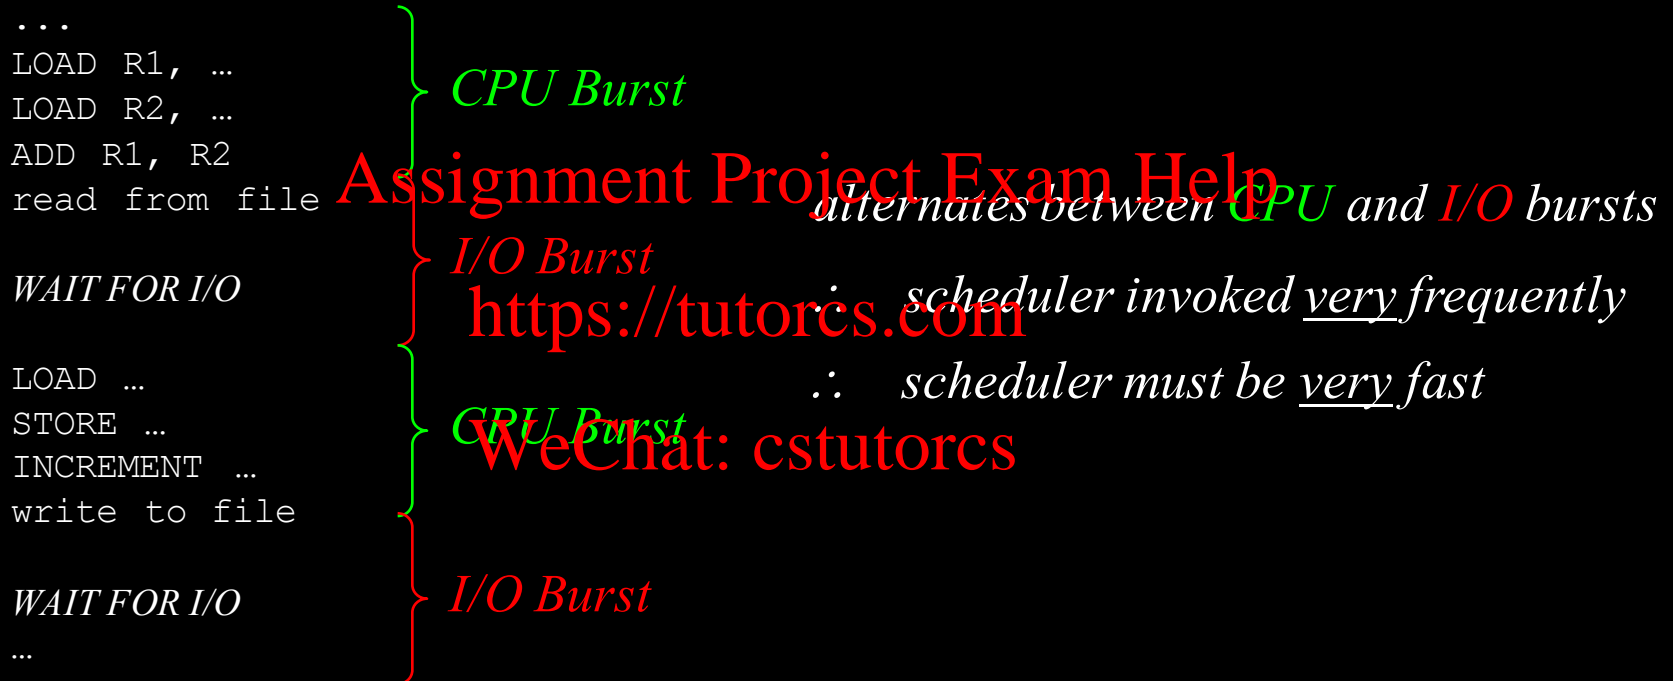
3. Examples

Basic Concepts

- One processor -> one process running at a time
 - All other processors wait for CPU to be free
 - Often running process does need CPU (e.g., I/O requests)
- Multiprogramming goal: avoid CPU idle times-> some process is using the CPU all the time
 - Keep multiple processes in memory at one time
 - When one process has to wait the CPU is allocated to another process
 - Scheduling algorithms: decide which process goes into the CPU

CPU-I/O Burst Cycle

A Typical Process Execution Flow

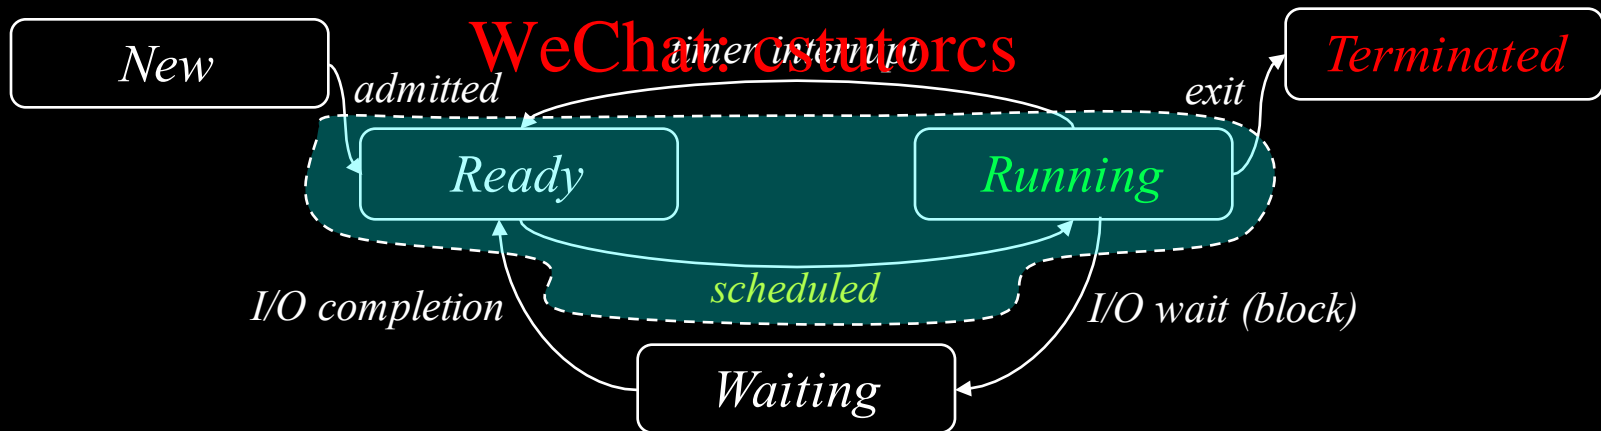


Distribution of CPU vs I/O bursts are important in selecting the right CPU scheduling approach

CPU Scheduler

CPU (Short-Term) Scheduler

- first example of resource manager (manages sharing of CPU)
- applicable to either **multiprocessing** or multithreading (we'll assume **former**)
- What? A kernel process that is always active (**daemon**)
- responsible for choosing process to run from ready queue



Agenda

1. Scheduling Overview

1. *First Example of Resource (CPU) Management (Sharing)*
- ➡ 2. *Non-Preemptive (N) vs. Preemptive Scheduling (P)*
3. *Metrics: Ways to Assess Effectiveness of Scheduling Policies*

Assignment Project Exam Help

2. Scheduling Policies

1. *First-Come-First-Served (FCFS) (N)*
2. *Shortest-Job-First (SJF) (N, P)*
3. *Priority (N, P)*
4. *Round-Robin (RR) (P)*
5. *Multilevel Queues (MLQ) (P) , Lottery*
6. *Real-time*

<https://tutorcs.com>

WeChat: cstutorcs

3. Examples

CPU Scheduling

When Does the Scheduler Get Invoked?

1. *when there's an interrupt*
 - *from timer: moves process from running to ready state*
 - *upon I/O interrupt: moves process from wait to ready state*

<https://tutorcs.com>

1. *when running process terminates*

WeChat: cstutorcs

2. *when running process issues an I/O request (block) and moves to wait state*

CPU Scheduling

Two Types of Scheduling:

Non-Preemptive: processes only give up CPU voluntarily

- simpler to implement (no timer interrupts)
- greedy or buggy process can starve others

Preemptive: processes also may be preempted by an interrupt

- adds complexity
- adds protection, better at ensuring fairness, as well as doing better in other scheduling metrics

CPU Scheduling Policies

1. First Come First Served

- *Non-preemptive*

2. Shortest Job First

- *a) Non-preemptive and b) Preemptive*

3. Priority Scheduling

- *a) Non-preemptive and b) Preemptive*

4. Round-Robin

- *Preemptive*

5. Multilevel Queues, Lottery

- *Preemptive*

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

How to compare these?

- *requires appropriate
scheduling metrics*

Agenda

1. Scheduling Overview

1. *First Example of Resource (CPU) Management (Sharing)*
2. *Non-Preemptive (N) vs. Preemptive Scheduling (P)*
- ➡ 3. *Metrics: Ways to Assess Effectiveness of Scheduling Policies*

Assignment Project Exam Help

2. Scheduling Policies

1. *First-Come-First-Served (FCFS) (N)*
2. *Shortest-Job-First (SJF) (N, P)*
3. *Priority (N, P)*
4. *Round-Robin (RR) (P)*
5. *Multilevel Queues (MLQ) (P) , Lottery*
6. *Real-time*

<https://tutorcs.com>

WeChat: cstutorcs

3. Examples

Scheduling Metrics

Typical Metrics for Comparing Scheduling Policies:

1. **Throughput**: measured in # processes completed / time unit
2. **Turnaround Time**: average (time of completion – time of creation)
 - related metric: **Waiting Time** (average time on ready queue)
 - **Waiting vs Turnaround**: no penalty for processes w/ long processing times
(Turnaround Time = Processing / Blocked Time + Waiting Time)
3. **Response Time**: time of 1st response is produced – time of creation
4. **Overhead**: time spent related to scheduling (e.g., context switch time)
5. **Fairness**: how much variation in waiting time
 - a) minimal
 - b) proportional (to processing time)
 - c) never infinite (aka: no starvation)

Can't Meet All Performance Goals With Any 1 Policy

Scheduling Performance Metrics

How Do We Measure Effectiveness of Scheduling Policies?

1. CPU Utilization

What percentage of time is the CPU doing useful work? (Good = High)

2. Throughput:

Number of processes that complete execution per time unit (Good = High)

<https://tutorcs.com>

3. Turnaround time:

Amount of time to execute a particular process. Includes time in ready and waiting queues (Good = Low)

WeChat: cstutorcs

4. Waiting time:

Average time process spends in Ready Queue (Good = Low)

5. Response time:

Average time before process produce first response after request (Good = Low)

Scheduling Metrics

Can't Meet All Performance Goals With 1 Policy

∴ Pick 1 or 2 that matter most

Typically Important Metrics:

1. Throughput / Turnaround time:
Throughput = n processes / sec \Rightarrow Turnaround = $1/n$ sec per process
<https://tutorcs.com>
2. Waiting time
3. Response time: (especially for highly interactive systems)
[WeChat: cstutorcs](https://tutorcs.com)

Also Care About:

1. Dispatch time:
Time it takes to choose next running process including schedule time + context switch time. If lengthy, effects effectiveness of scheduler
2. Fairness: Want to avoid process *starvation*

Agenda

1. Scheduling Overview

1. *First Example of Resource (CPU) Management (Sharing)*
2. *Non-Preemptive (N) vs. Preemptive Scheduling (P)*
3. *Metrics: Ways to Assess Effectiveness of Scheduling Policies*

Assignment Project Exam Help

2. Scheduling Policies

<https://tutorcs.com>



1. *First-Come-First-Served (FCFS) (N)*
2. *Shortest-Job-First (SJF) (N, P)*
3. *Priority (N, P)*
4. *Round-Robin (RR) (P)*
5. *Multilevel Queues (MLQ) (P)*
6. *Multiprocessor, energy aware, real time, overload*

WeChat: cstutorcs

3. Examples

Assignment Project Exam Help

Schedule Policy #1:

<https://tutores.com>

FCFS: First Come First Served

WeChat: cstutores

1. First-Come First-Serve (FCFS)

Non-Preemptive Policy

Example

Process

CPU Burst Time

P_1

24

P_2

3

P_3

3

Assignment Project Exam Help

<https://tutorcs.com>

Suppose processes arrive in order: $P_1; P_2; P_3$



(*Gantt Chart*)

Waiting Time

$P_1 : 0$

$P_2 : 24$

$P_3 : 27$

} *Average Wait Time = 17*

1. First-Come First-Serve (FCFS)

Non-Preemptive Policy

Example

<u>Process</u>	<u>CPU Burst Time</u>
----------------	-----------------------

P_1	24
-------	----

P_2	3
-------	---

P_3	3
-------	---

Assignment Project Exam Help

Suppose processes arrive in order: P_2, P_3, P_1



Waiting Time

$P_1 : 6$

$P_2 : 0$

$P_3 : 3$

Average Wait Time = 3

*Waiting time
determined by
order of arrival*

1. First-Come First-Serve (FCFS)

Non-Preemptive Policy

+: *Easy to implement (ready queue = scheduling queue)*

No starvation

Assignment Project Exam Help

–: *Convoy Effect: Order of arrival determines performance*

<https://tutorcs.com>

e.g., many processes might end up waiting for a big one to get off the CPU

WeChat: cstutorcs

(Food for thought: maybe shortest process should go first instead?)

CPU Scheduling Policies Summary

<i>Policy</i>	<i>Throughput</i>	<i>Waiting</i>	<i>Response</i>	<i>Fairness</i>	<i>Overhead</i>	<i>Comments</i>
<i>FCFS</i>	✗	✗	✗	✓	✓	Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Schedule Policy #2:
Assignment Project Exam Help
SJF: Shortest Job First
<https://tutorcs.com>

WeChat: cstutorcs

2a. Shortest-Job-First (SJF)

Non-Preemptive Policy

Idea: Rank processes by CPU time requests.
Optimizes (minimizes) average waiting time

Example

<u>Process</u>	<u>CPU Burst</u>
P ₁	24
P ₂	3
P ₃	3
P ₄	7

<https://tutorcs.com>
WeChat: cstutorcs



$$\text{Average wait} = (0 + 3 + 6 + 13) / 4 = 5.5$$


2a. Shortest-Job-First (SJF)

Q: What if processes don't arrive at same time?

A: Whenever process finishes, choose next process with shortest CPU burst

Example

Process	Arrival Time	CPU Burst
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4



Average wait time: ?


2a. Shortest-Job-First (SJF)

Q: What if processes don't arrive at same time?

A: *Whenever process finishes, choose next process with shortest CPU burst*

Example

Process	Arrival Time	CPU Burst
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4



Average wait time: $(0+6+3+7) / 4 = 4$

Preemptive version of SJF does even better ...

2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4


2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4



$P_1 = 5, P_2 = 4$


2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs

Example	Process	Arrival Time	CPU Burst
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs

Example	Process	Arrival Time	CPU Burst
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

Example	Process	Arrival Time	CPU Burst
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4



2b. Preemptive SJF

Example:

<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4



Average wait time:

$$\left. \begin{array}{l} P_1 : 9 \\ P_2 : 1 \\ P_3 : 0 \\ P_4 : 2 \end{array} \right\} \begin{array}{l} \text{Average} = 3 \\ \text{(Compare to non-preemptive SJF average} = 4) \end{array}$$