

COSI 131a: Operating Systems

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs (2)

Chapter 6

Agenda

1. Scheduling Overview

1. *First Example of Resource (CPU) Management (Sharing)*
2. *Non-Preemptive (N) vs. Preemptive Scheduling (P)*
3. *Metrics: Ways to Assess Effectiveness of Scheduling Policies*

Assignment Project Exam Help

2. Scheduling Policies

<https://tutorcs.com>

1. *First-Come-First-Served (FCFS) (N)*



2. *Shortest-Job-First (SJF) (N, P)*

WeChat: cstutorcs

3. *Priority (N, P)*

4. *Round-Robin (RR) (P)*

5. *Multilevel Queues (MLQ) (P) , Lottery*

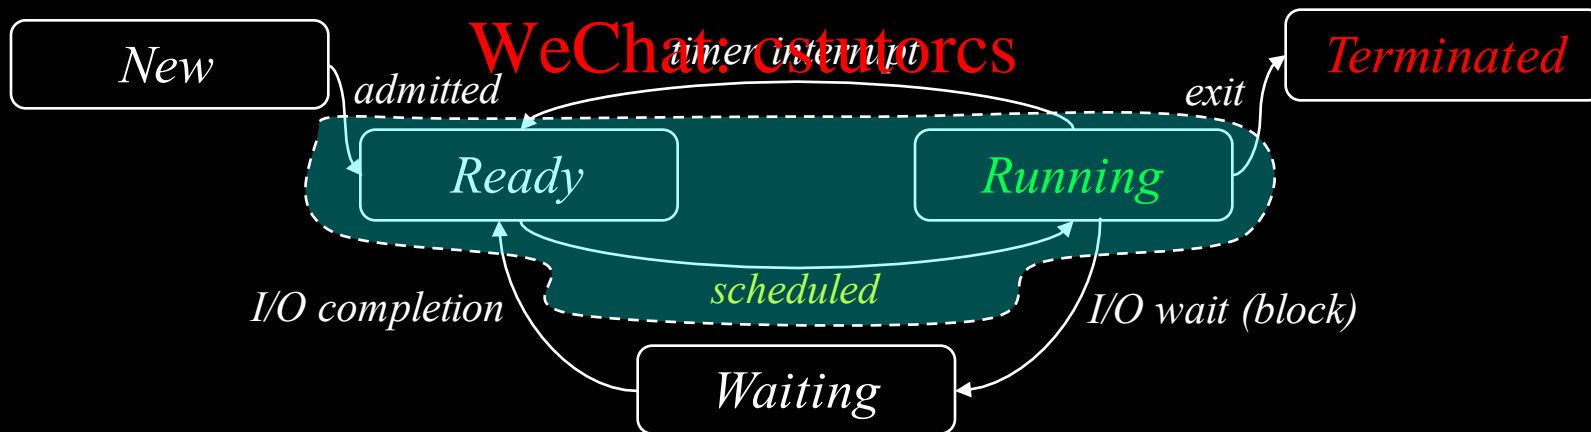
6. *Real-time*

3. Examples

Review: CPU Scheduler

CPU (Short-Term) Scheduler

- first example of resource manager (manages sharing of CPU)
- applicable to either **multiprocessing** or multithreading (we'll assume **former**)
- What? A kernel process that is always active (**daemon**)
- responsible for choosing process to run from ready queue



Review: CPU Scheduling

Two Types of Scheduling:

Non-Preemptive: processes only give up CPU voluntarily

- simpler to implement (no timer interrupts)
- greedy or buggy process can starve others

Preemptive: processes also may be preempted by an interrupt

- adds complexity
- adds protection, better at ensuring fairness, as well as doing better in other scheduling metrics

Review: Scheduling Metrics

Can't Meet All Performance Goals With 1 Policy

∴ Pick 1 or 2 that matter most

Typically Important Metrics:

1. Throughput / Turnaround time:
Throughput = n processes / sec \Rightarrow Turnaround = $1/n$ sec per process
<https://tutorcs.com>
2. Waiting time
3. Response time: (especially for highly interactive systems)
[WeChat: cstutorcs](https://tutorcs.com)

Also Care About:

1. Dispatch time:
Time it takes to choose next running process including schedule time + context switch time. If lengthy, effects effectiveness of scheduler
2. Fairness: Want to avoid process *starvation*

Review: First-Come First-Serve

Non-Preemptive Policy

Example

Process

CPU Burst Time

P_1

24

P_2

3

P_3

3

Assignment Project Exam Help

<https://tutorcs.com>

Suppose processes arrive in order: $P_1; P_2; P_3$



(*Gantt Chart*)

Waiting Time

$P_1 : 0$

$P_2 : 24$

$P_3 : 27$

} *Average Wait Time = 17*

CPU Scheduling Policies Summary

<i>Policy</i>	<i>Throughput</i>	<i>Waiting</i>	<i>Response</i>	<i>Fairness</i>	<i>Overhead</i>	<i>Comments</i>
<i>FCFS</i>	✗	✗	✗	✓	✓	Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Schedule Policy #2:
Assignment Project Exam Help
SJF: Shortest Job First
<https://tutorcs.com>

WeChat: cstutorcs

2a. Shortest-Job-First (SJF)

Non-Preemptive Policy

Idea: Rank processes by CPU time requests.
Optimizes (minimizes) average waiting time

Example

<u>Process</u>	<u>CPU Burst</u>
P ₁	24
P ₂	3
P ₃	3
P ₄	7

<https://tutorcs.com>
WeChat: cstutorcs



$$\text{Average wait} = (0 + 3 + 6 + 13) / 4 = 5.5$$


2a. Shortest-Job-First (SJF)

Q: What if processes don't arrive at same time?

A: Whenever process finishes, choose next process with shortest CPU burst

Example

Process	Arrival Time	CPU Burst
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4



Average wait time: ?


2a. Shortest-Job-First (SJF)

Q: What if processes don't arrive at same time?

A: Whenever process finishes, choose next process with shortest CPU burst

Example

Process	Arrival Time	CPU Burst
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4



Average wait time: $(0+6+3+7) / 4 = 4$

Preemptive version of SJF does even better ...

2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4


2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs

Example	Process	Arrival Time	CPU Burst
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4



P₁ = 5, P₂ = 4


2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs

Example	Process	Arrival Time	CPU Burst
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs

<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs

Example	Process	Arrival Time	CPU Burst
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4




2b. Preemptive SJF

Preemptive Version of SJF

Idea: *Currently running process can be preempted if new process arrives with shorter remaining CPU burst*

Assignment Project Exam Help
<https://tutorcs.com>
 WeChat: cstutorcs


<i>Example</i>	<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
	P ₁	0.0	7
	P ₂	2.0	4
	P ₃	4.0	1
	P ₄	5.0	4



2b. Preemptive SJF

Example:

<u>Process</u>	<u>Arrival Time</u>	<u>CPU Burst</u>
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4



Average wait time:

$$\left. \begin{array}{l} P_1 : 9 \\ P_2 : 1 \\ P_3 : 0 \\ P_4 : 2 \end{array} \right\} \begin{array}{l} \text{Average} = 3 \\ \text{(Compare to non-preemptive SJF average} = 4) \end{array}$$

2a, 2b: SJF

Issue: Predicting how much is next CPU burst

Exponential Averaging

Uses “history” of previous CPU bursts to predict length of next CPU burst

t_n : Actual length of n th CPU burst

τ_{n+1} : Predicted length of next CPU burst

α : “Weight”:

→ How much emphasis to put on recent past

→ $0 \leq \alpha \leq 1$

Policies # 2a,2b: SJF (cont.)

Exponential Averaging

$$\tau_{n+1} = (\alpha \times t_n) + ((1 - \alpha) \times \tau_n)$$

τ_{n+1} : *next guess*

τ_n : *previous guess (history)*


t_n : *previous actual burst time (most recent information)*

α : *weighting factor ($0 \leq \alpha \leq 1$)*

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

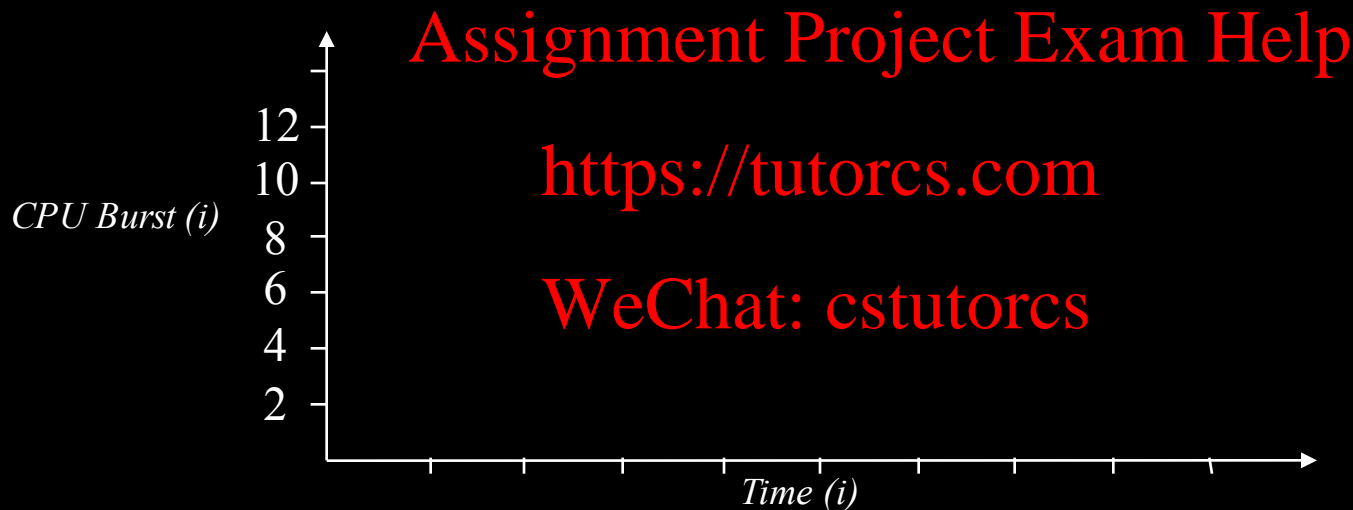
Example: $\alpha = 0.5$ $\tau_{n+1} = (\alpha \times t_n) + ((1 - \alpha) \times \tau_n)$



Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

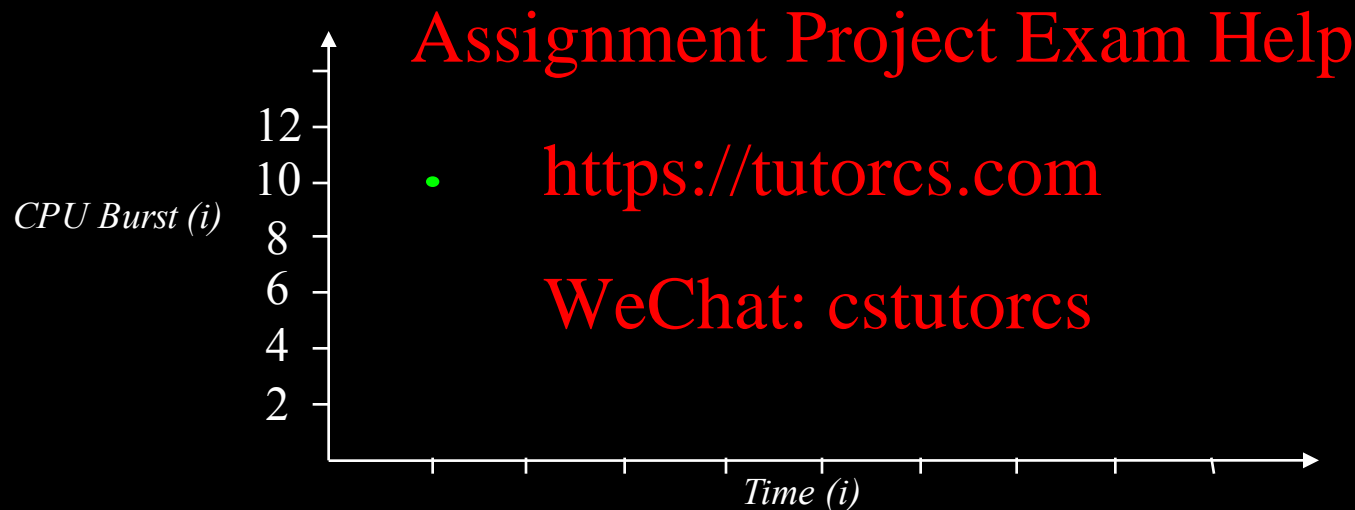


Actual (t_i):	
Guess (τ_i):	

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

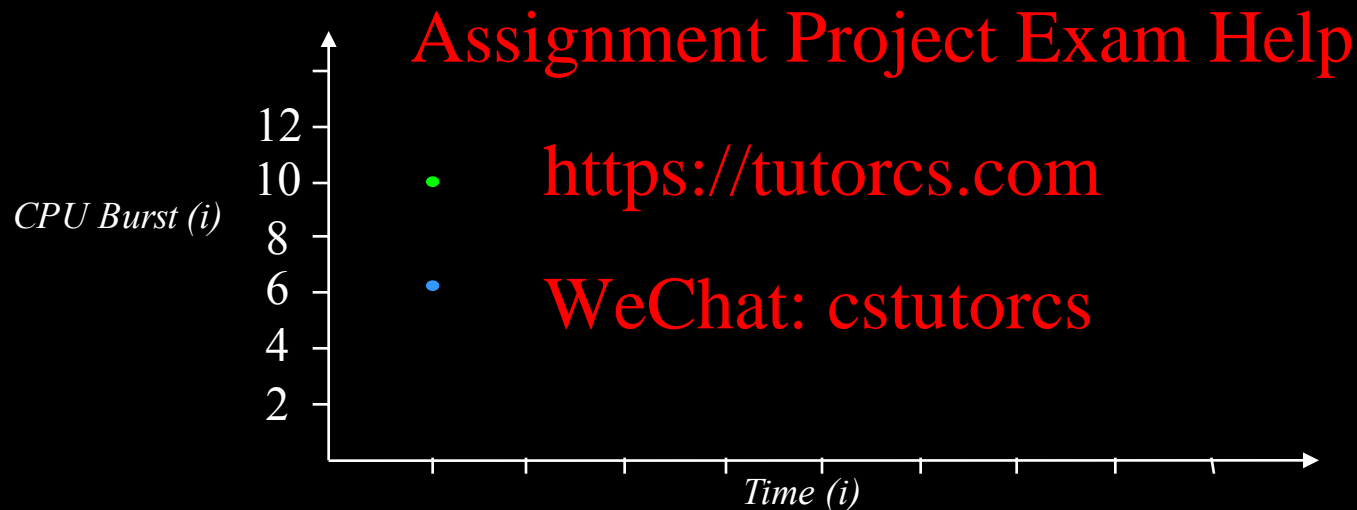


Actual (t_i):	
Guess (τ_i):	10

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

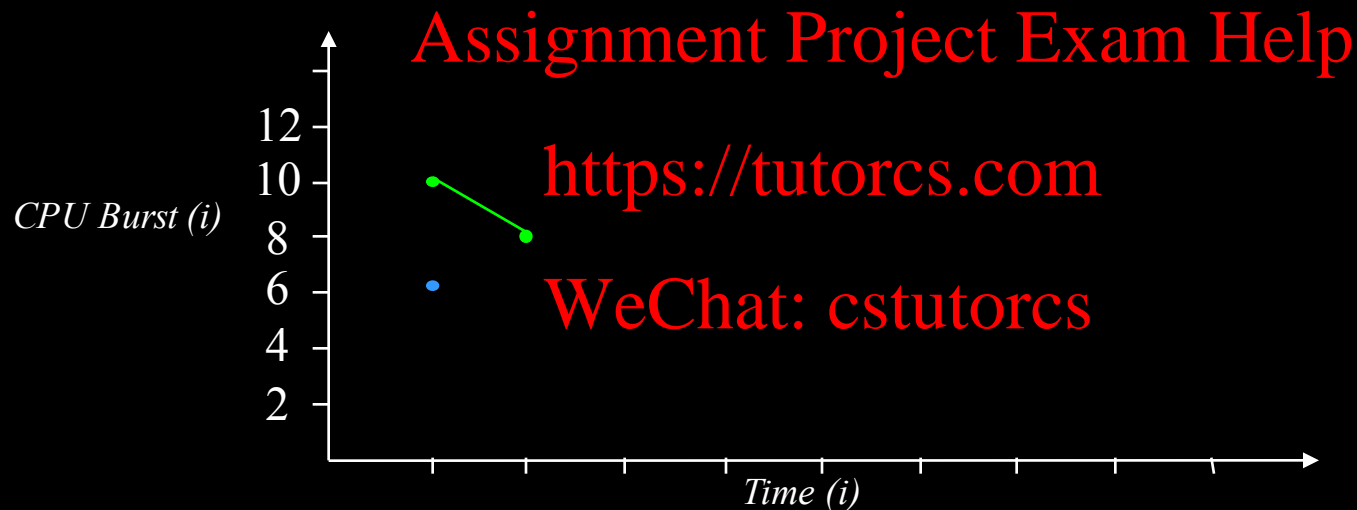


Actual (t_i):	6
Guess (τ_i):	10

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

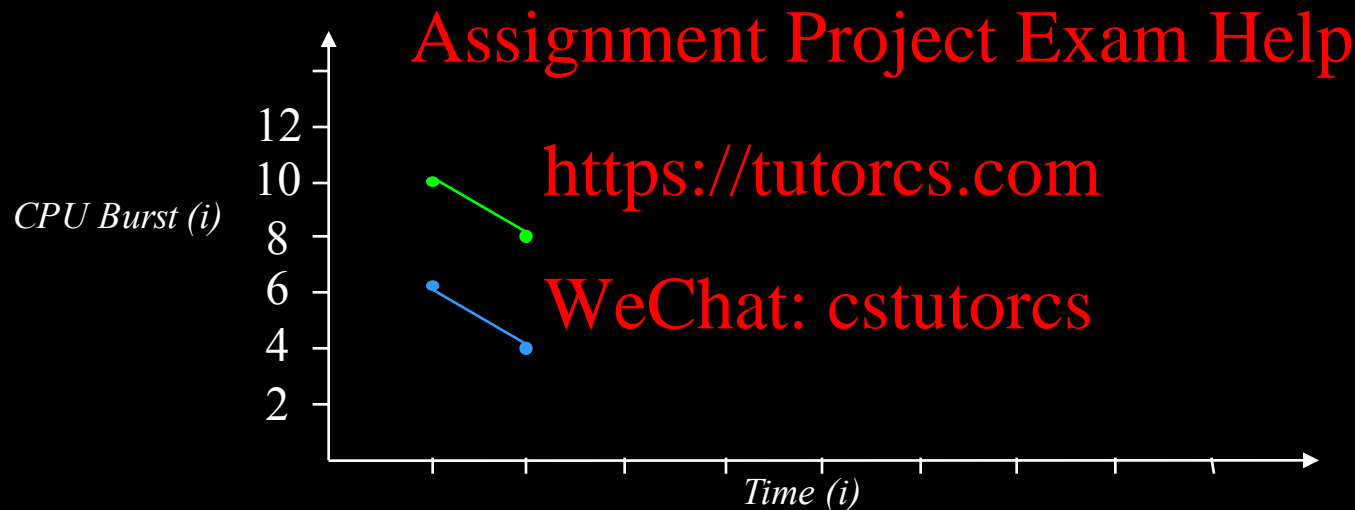


Actual (t_i):	6
Guess (τ_i):	10 8

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

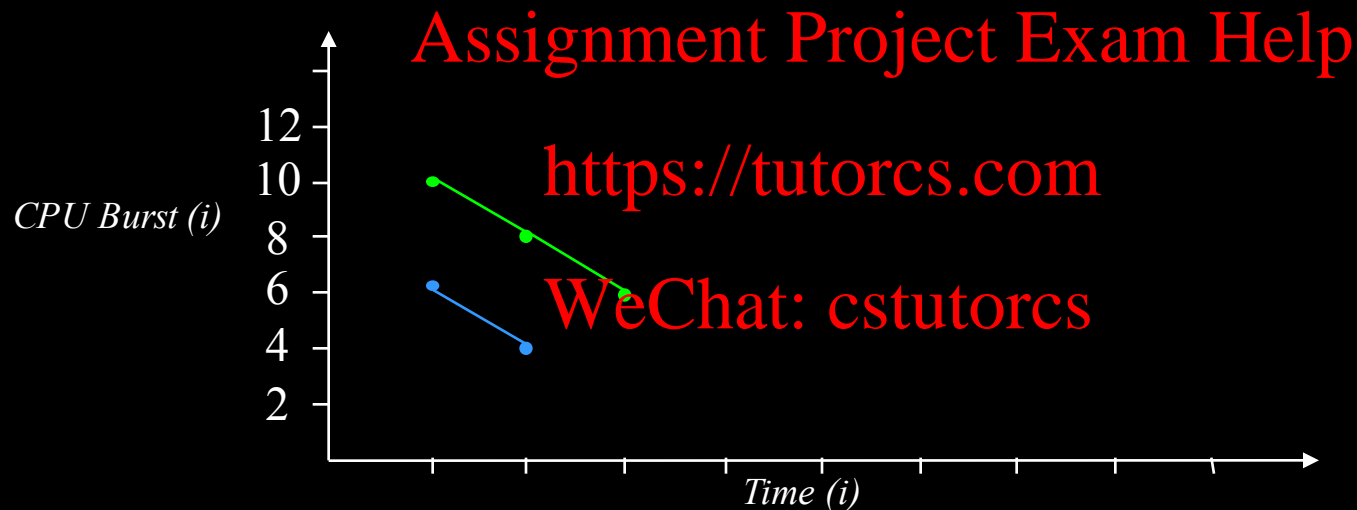


Actual (t_i):	6	4
Guess (τ_i):	10	8

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

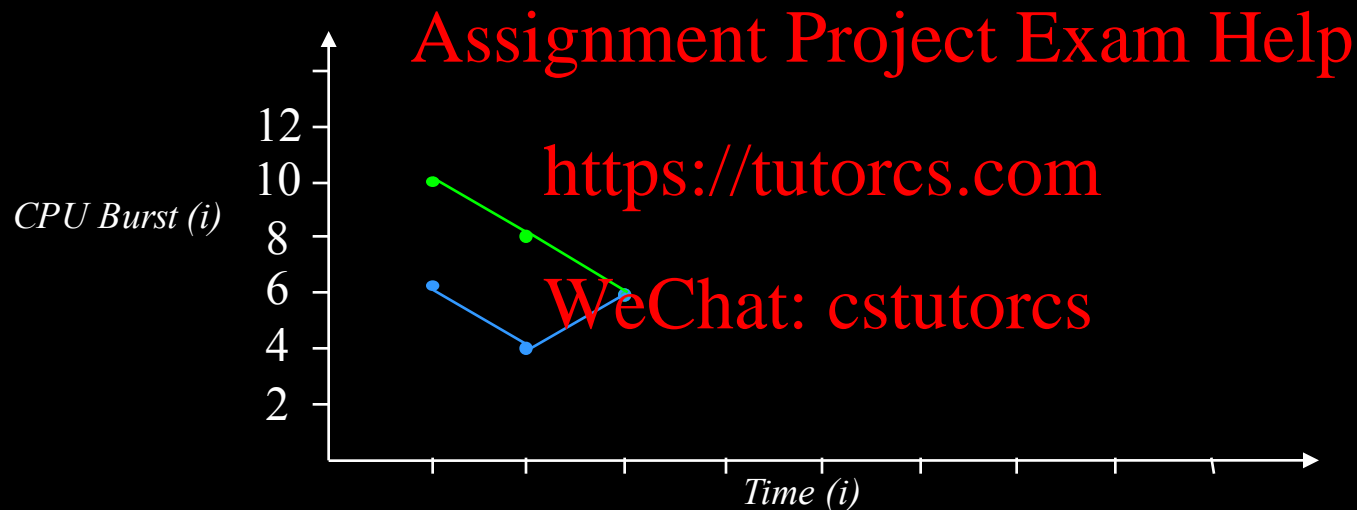


Actual (t_i):	6	4	
Guess (τ_i):	10	8	6

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

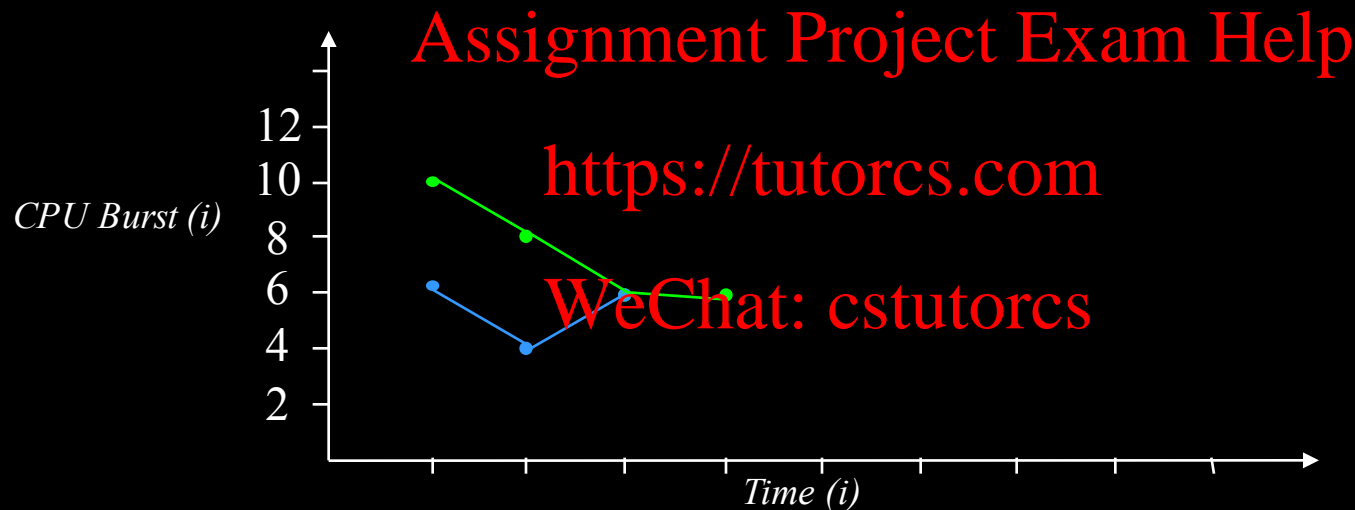


Actual (t _i):	6	4	6
Guess (τ _i):	10	8	6

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

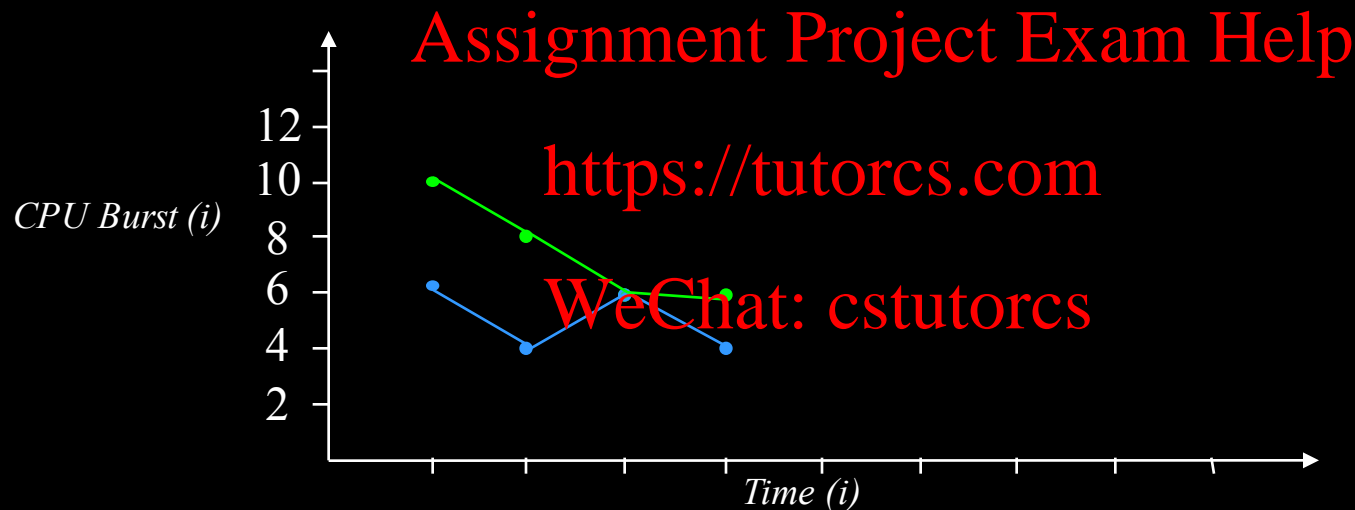


Actual (t_i):	6	4	6	
Guess (τ_i):	10	8	6	6

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

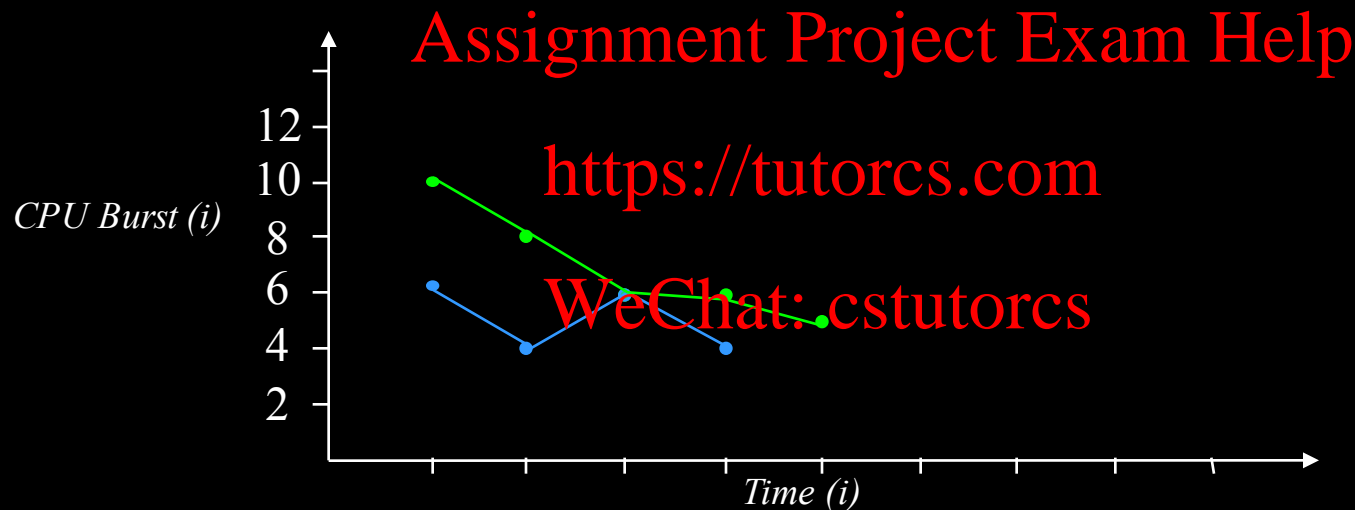


Actual (t_i):	6	4	6	4
Guess (τ_i):	10	8	6	6

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$




Actual (t_i):	6	4	6	4
Guess (τ_i):	10	8	6	5

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

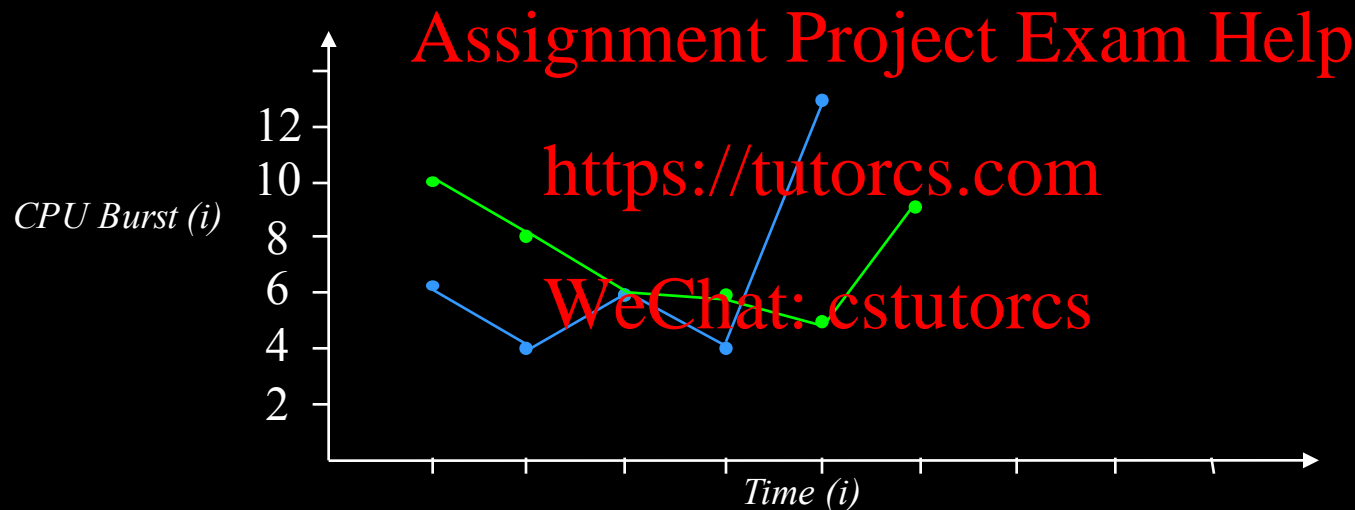


Actual (t_i):	6	4	6	4	13
Guess (τ_i):	10	8	6	6	5

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

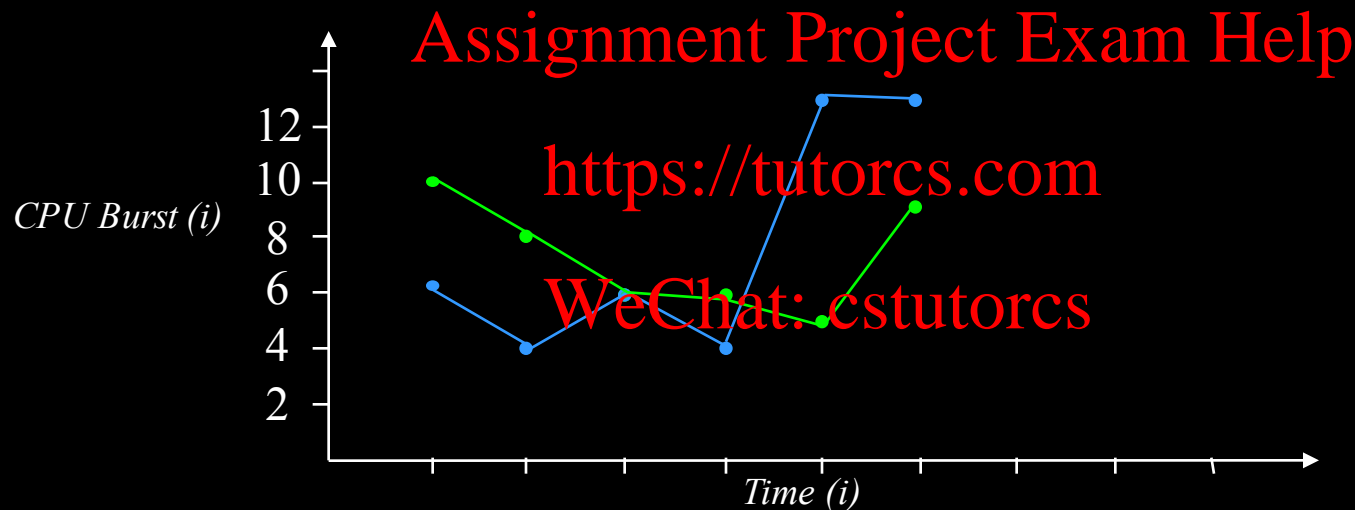


Actual (t_i):	6	4	6	4	13	
Guess (τ_i):	10	8	6	6	5	9

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

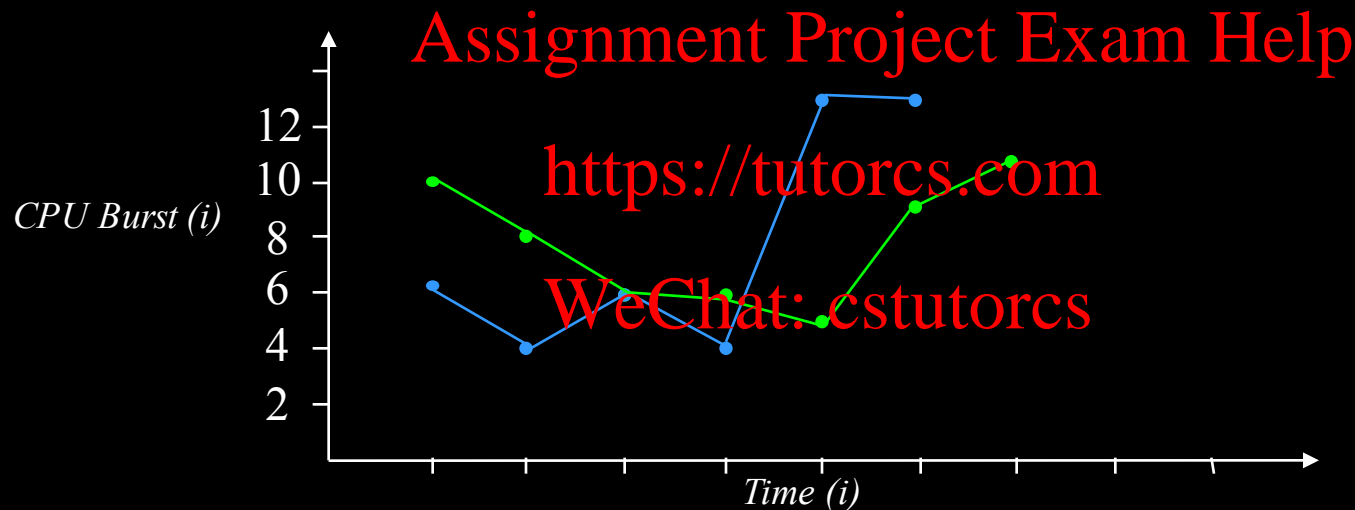


Actual (t_i):	6	4	6	4	13	13
Guess (τ_i):	10	8	6	6	5	9

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

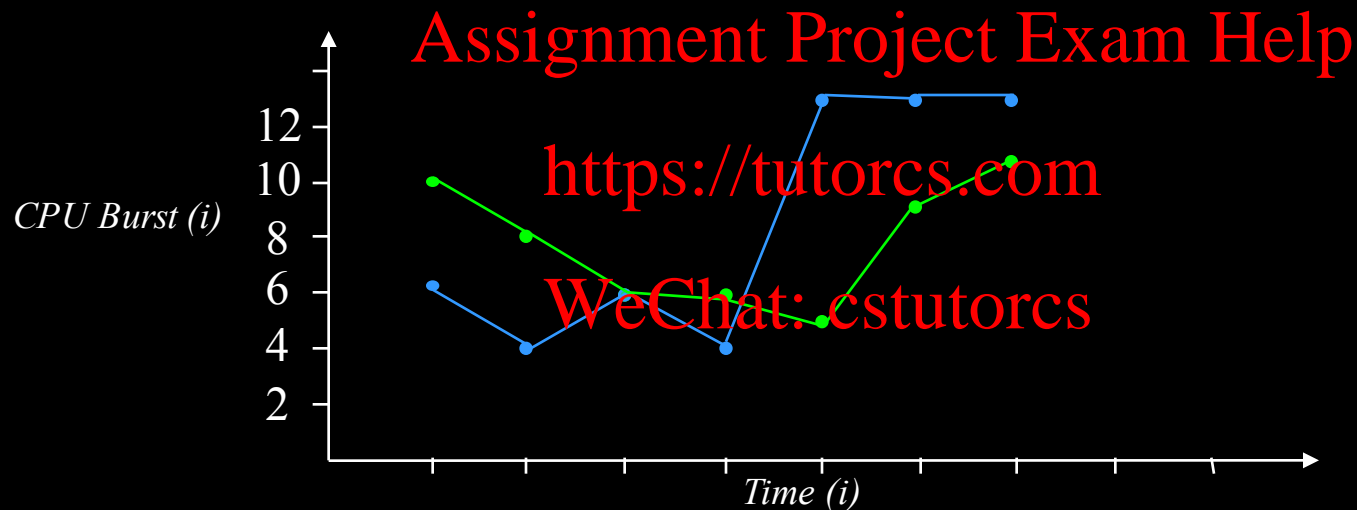


Actual (t_i):	6	4	6	4	13	13	
Guess (τ_i):	10	8	6	6	5	9	11

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$

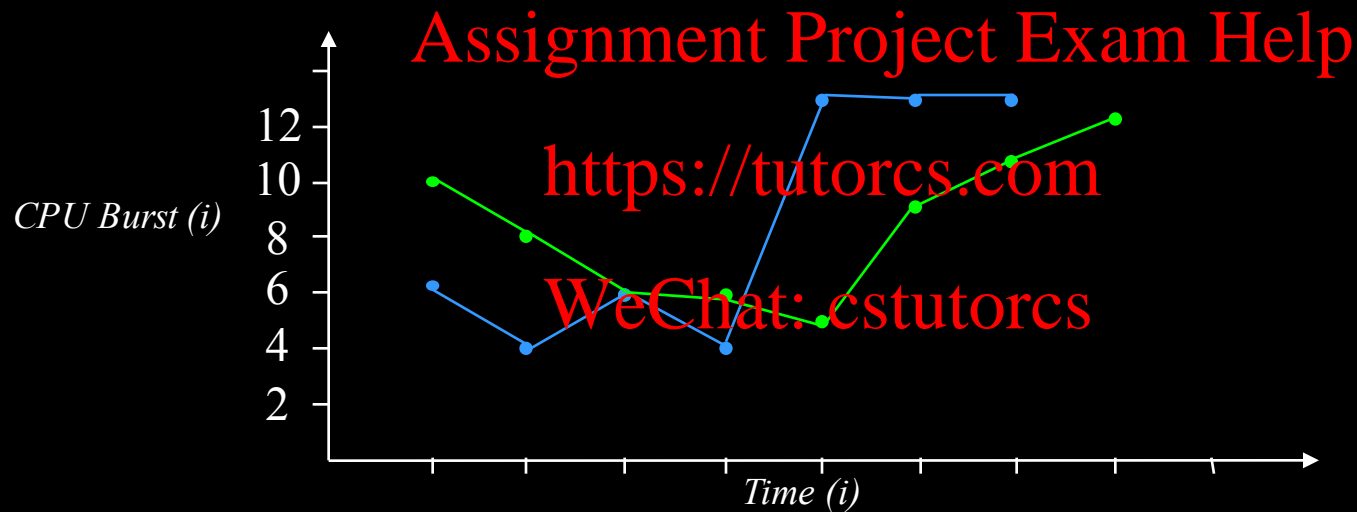


Actual (t_i):	6	4	6	4	13	13	13
Guess (τ_i):	10	8	6	6	5	9	11

Policies # 2a,2b: SJF (cont.)

Exponential Averaging (cont.)

Example: $\alpha = 0.5$ $\tau_{n+1} = (0.5 \times t_n) + (0.5 \times \tau_n)$



Actual (t_i):	6	4	6	4	13	13	13	
Guess (τ_i):	10	8	6	6	5	9	11	12

Policies # 2a,2b: SJF (cont.)

Exponential Average:

$$\tau_{n+1} = (\alpha \times t_n) + ((1 - \alpha) \times \tau_n)$$

Choosing α :

1. $\alpha = 0 \Rightarrow \tau_{n+1} = \tau_n$
 • Recent history doesn't count

2. $\alpha = 1 \Rightarrow \tau_{n+1} = t_n$
<https://tutorcs.com>

• Only most recent burst counts

WeChat: cstutorcs

3. In general: $\tau_{n+1} = (\alpha \times t_n) +$

Weights decrease
 since $(1 - \alpha) \leq 1$

$$\begin{aligned} & \frac{(1 - \alpha) \times \alpha}{1} t_{n-1} + \\ & \frac{(1 - \alpha)^2 \times \alpha}{2} t_{n-2} + \\ & \frac{(1 - \alpha)^3 \times \alpha}{3} t_{n-3} + \\ & \dots \\ & \frac{(1 - \alpha)^n \times \alpha}{n} t_0 + \end{aligned}$$

Policies # 2a,2b: SJF (cont.)

+: *Approximates optimal schedule for*

- *Average wait time*
- *Average turnaround time*

Assignment Project Exam Help

—: *Fairness*

<https://tutorcs.com>

- *Starves processes with expected long CPU bursts*

WeChat: cstutorcs

CPU Scheduling Policies Summary

<i>Policy</i>	<i>Throughput</i>	<i>Waiting</i>	<i>Response</i>	<i>Fairness</i>	<i>Overhead</i>	<i>Comments</i>
<i>FCFS</i>	✗	✗	✗	✓	✓	+ <i>easy implementation</i> - <i>convoy effect</i>

<https://tutorcs.com>

WeChat: cstutorcs

CPU Scheduling Policies Summary

<i>Policy</i>	<i>Throughput</i>	<i>Waiting</i>	<i>Response</i>	<i>Fairness</i>	<i>Overhead</i>	<i>Comments</i>
<i>FCFS</i>	✗	✗	✗	✓	✓	+ <i>easy implementation</i> - <i>convoy effect</i>
<i>SJF</i>	✓	✓	✗	✗		+ <i>waiting turnaround time</i> -- <i>starvation</i> <i>hard to predict CPU time</i>

Assignment Project Exam Help


Schedule Policy #3: Priority-Based

<https://tutorcs.com>

WeChat: cstutorcs

Scheduling the Ready Queue

Processes Not Always Running



New:

Process is being created

Running:

Instructions are being executed

Waiting:

Process is waiting for some event to occur

Ready:

Process is waiting to be assigned the CPU

Terminated:

Process has finished execution

Policies # 3a,3b: Priority

Idea:

- *Priority Number (Integer) Associated With Each Process*
- *CPU Allocated to Process With Highest Priority (Lowest Priority Number)...*

<https://tutorcs.com>

3a. Non-preemptive:

Whenever running process blocks or terminates

3b. Preemptive:

Whenever new process arrives or running process blocks or terminates


Example ...

Policies # 3a,3b: Priority

Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst</u>	<u>Priority</u>
P ₁	0.0	7	2
P ₂	2.0	4	0
P ₃	4.0	1	3
P ₄	5.0	4	1

Assignment Project Exam Help



Q: How does priority scheduling generalize SJF scheduling?

A: SJF = Priority scheduling where shortest CPU burst = highest priority

Policies # 3a,3b: Priority

Assessing Priority Scheduling

+:

Reflects relative importance of processes (e.g.: kernel > user)

-:

Assignment Project Exam Help

Turnaround, wait times (process burst length ignored)

Starvation of low-priority processes

→ Can counter with “aging”
(process gets increasing priority the more it waits)

CPU Scheduling Policies Summary

<i>Policy</i>	<i>Throughput</i>	<i>Waiting</i>	<i>Response</i>	<i>Fairness</i>	<i>Overhead</i>	<i>Comments</i>
<i>FCFS</i>	✗	✗	✗	✓	✓	+ <i>easy implementation</i> - <i>convoy effect</i>
<i>SJF</i>	✓	✓		✗		+ <i>waiting turnaround time</i> -- <i>starvation</i> <i>hard to predict CPU time</i>

WeChat: cstutores

CPU Scheduling Policies Summary

<i>Policy</i>	<i>Throughput</i>	<i>Waiting</i>	<i>Response</i>	<i>Fairness</i>	<i>Overhead</i>	<i>Comments</i>
<i>FCFS</i>	✗	✗	✗	✓	✓	+ <i>easy implementation</i> - <i>convoy effect</i>
<i>SJF</i>	✓	✓	✗	✗		+ <i>waiting turnaround time</i> -- <i>starvation</i> <i>hard to predict CPU time</i>
<i>Priority</i>				✗		

Assignment Project Exam Help

Schedule Policy #4: Round-Robin

<https://tutorcs.com>

WeChat: cstutorcs

Policy #4: Round-Robin

Idea:

- Each process gets small unit of CPU time (*time quantum*).
- After time quantum has elapsed, process is preempted and added to end of ready queue
- Preemptive only: Process switch based on timer interrupts

Goal: Fairness and Response Time

Suppose Time Quantum = q , & n Processes in Ready Queue

Each process gets $1/n$ of CPU time in chunks of q

No process waits more than $(n - 1) \times q$ time units

Policy # 4: Round-Robin

Example: *Time Quantum = 20*

Process	CPU Burst
P ₁	0
P ₂	0
P ₃	0
P ₄	0

Assignment Project Exam Help

<https://tutorcs.com>



Waiting time:

$$P_1 : 17 + 20 + 20 + 20 + 4 = 81$$

$$P_2 : 20$$

$$P_3 : 20 + 17 + 20 + 20 + 4 + 13 = 94$$

$$P_4 : 20 + 17 + 20 + 20 + 20 = 97$$

Average: 73

Policy #4: Round-Robin

+: *Fairness*

Response time

Assignment Project Exam Help


–: *Turnaround time* <https://tutorcs.com>

Cost of context switches (Must choose quantum carefully)

WeChat: [cstutorcs](#)

Policy #4: Round-Robin

Choosing a Time Quantum, q :



1. Large $q \Rightarrow$ few context switches.
 $q > \text{all CPU bursts} = \text{FIFO}$
2. Small $q \Rightarrow$ many context switches.
 Too small \Rightarrow overhead can compromise performance

Policy #4: Round-Robin

Choosing a Time Quantum, q

Should be large with respect to context switch time

$$\frac{\text{context-switch-time}}{\text{context-switch-time} + q} = \% \text{ of CPU spent switching}$$

Should also consider process performance (e.g.: Average turnaround time)

Q: Given:

<u>Process</u>	<u>CPU Burst</u>
P ₁	6
P ₂	3
P ₃	1
P ₄	7

What is Average Turnaround Time Assuming $q = 2$?

Policy #4: Round-Robin

$q = 2:$

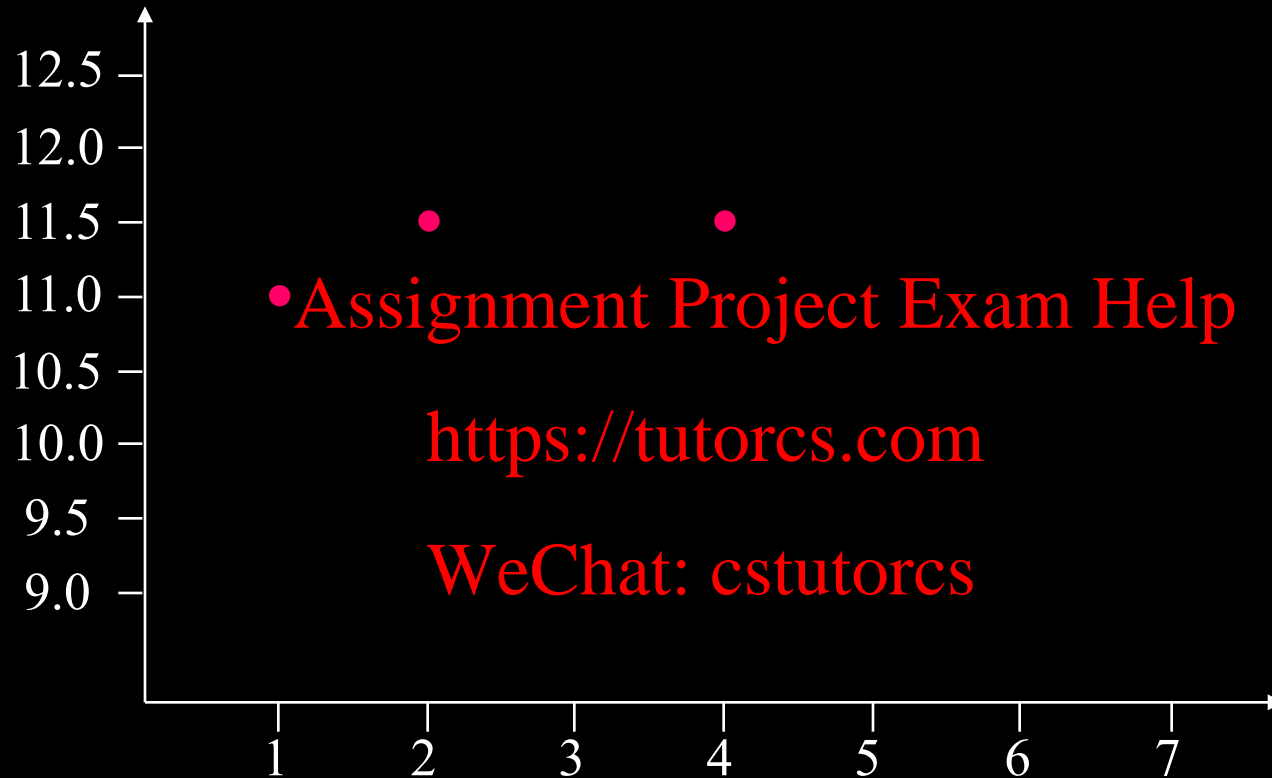
<u>Process</u>	<u>CPU Burst</u>
P ₁	0
P ₂	0
P ₃	0
P ₄	0



Turnaround time:

P ₁ : 14	} <i>Average = 11.5</i>
P ₂ : 10	
P ₃ : 5	
P ₄ : 17	

Policy #4: Round-Robin



Q: Determine the Average Turnaround Time Assuming $q = \dots$

- a) 3 b) 5 c) 6 d) 7

Policy #4: Round-Robin

$q = 3, 5, 6, 7$:

<u>Process</u>	<u>CPU Burst</u>
P ₁	6
P ₂	5
P ₃	1
P ₄	7

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs
Turnaround time?

Policy #4: Round-Robin

$q = 3$:

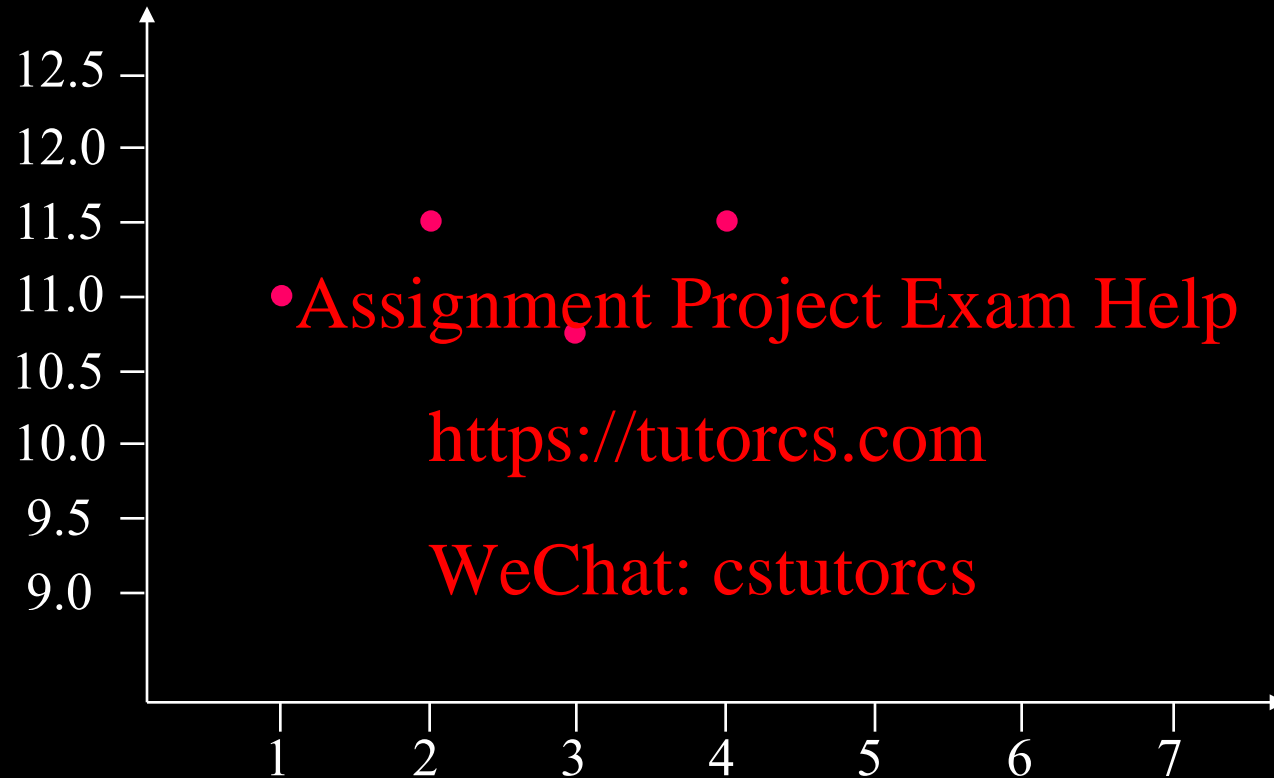
<u>Process</u>	<u>CPU Burst</u>
P ₁	6
P ₂	3
P ₃	1
P ₄	7



Turnaround time:

$$\left. \begin{array}{l} P_1 : 13 \\ P_2 : 6 \\ P_3 : 7 \\ P_4 : 17 \end{array} \right\} \text{Average} = 10.75$$


Policy #4: Round-Robin



Policy #4: Round-Robin

$q = 5$

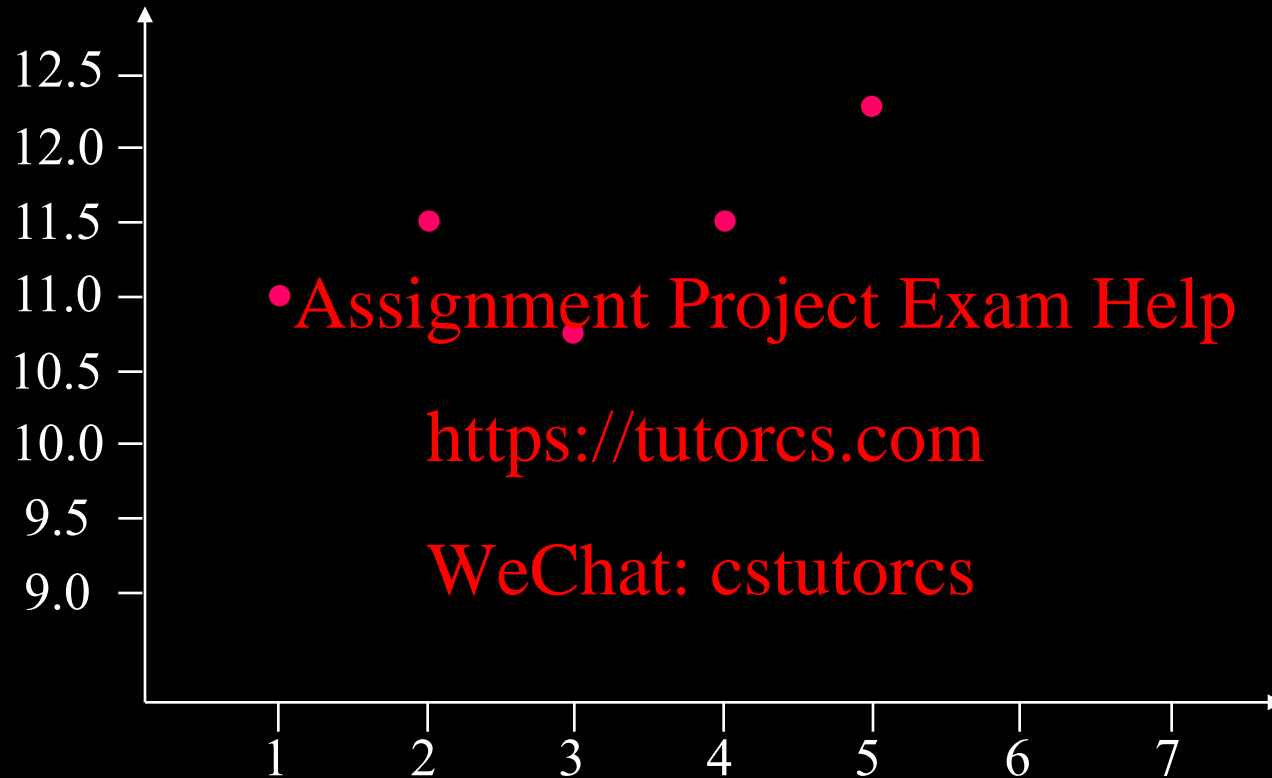
Process	CPU Burst
P ₁	6
P ₂	3
P ₃	1
P ₄	7



Turnaround time:

$$\left. \begin{array}{l} P_1 : 15 \\ P_2 : 8 \\ P_3 : 9 \\ P_4 : 17 \end{array} \right\} \text{Average} = 12.25$$


Policy #4: Round-Robin



Policy #4: Round-Robin

$q = 6,$
 $q = 7$


<u>Process</u>	<u>CPU Burst</u>
P ₁	6
P ₂	3
P ₃	1
P ₄	7



Turnaround time:

P ₁ : 6	} <i>Average = 10.5</i>
P ₂ : 9	
P ₃ : 10	
P ₄ : 17	

Policy #4: Round-Robin



CPU Scheduling Policies Summary

<i>Policy</i>	<i>Throughput</i>	<i>Waiting</i>	<i>Response</i>	<i>Fairness</i>	<i>Overhead</i>	<i>Comments</i>
<i>FCFS</i>	✗	✗	✗	✓	✓	+ <i>easy implementation</i> - <i>convoy effect</i>
<i>SJF</i>	✓✓	✓✓	✗	✗		+ <i>waiting turnaround time</i> -- <i>starvation</i> <i>hard to predict CPU time</i>
<i>Priority</i>				✗		+ <i>importance considered</i> -- <i>poor metrics for low priority</i>

CPU Scheduling Policies Summary

<i>Policy</i>	<i>Throughput</i>	<i>Waiting</i>	<i>Response</i>	<i>Fairness</i>	<i>Overhead</i>	<i>Comments</i>
<i>FCFS</i>	✗	✗	✗	✓	✓	+ <i>easy implementation</i> - <i>convoy effect</i>
<i>SJF</i>	✓	✓	✗	✗		+ <i>waiting turnaround time</i> -- <i>starvation</i> <i>hard to predict CPU time</i>
<i>Priority</i>				✗		+ <i>importance considered</i> -- <i>poor metrics for low priority</i>
<i>Round-Robin</i>			✓	✓	✗	

Revise: Throughput vs Turnaround

Analogy: Plane vs Train for Shipping Tables (MA-CA)

Plane: 6 hours of travel each way, 48 tables per trip

Train: 2 days of travel each way, $60 \times 96 = 5760$ tables per trip

Processing: Can assemble 60 tables / hour

Turnaround Time = max time to "process" a table

<https://tutorcs.com>

lowest *Plane: 12 (wait) + 6 (travel) + 0.8 hrs (assembling) ~ 19 hrs*

Train: 96 (wait) + 48 (travel) + 96 (processing) = 10 days

WeChat: cstutorcs

Throughput = number of tables processed / day

	Day 1	Day 2	Day 3	Day 4	Day 5	Day 6	Throughput
<i>Plane</i>	96	96	96	96	96	96	96 / day
<i>Train</i>	0	0	1440	1440	1440	1440	~1440 / day

Assignment Project Exam Help

Schedule Policy #5: Multilevel Queues

<https://tutorcs.com>

WeChat: cstutorcs

Policy #5: Multilevel Queues

Idea:

- *Ready queue is partitioned into separate queues, each with own scheduling algorithm*

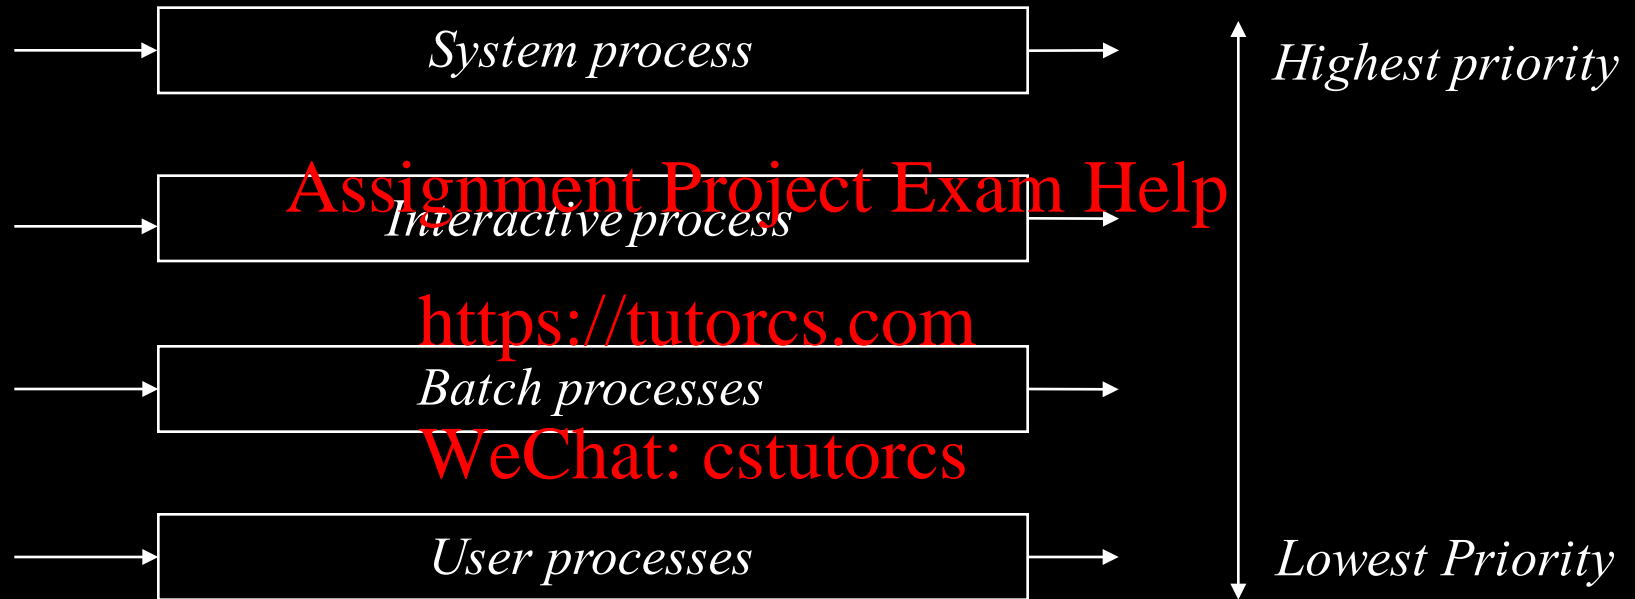
→ e.g. Queue 1 (for interactive processes): RR
Queue 2 (for others): FCFS

Requires Interqueue Scheduling too (Scheduling of Queues)

1. Fixed Priority:
Serve all from high-priority queue before servicing lower-priority queue (–: starvation)
2. Time-slice:
Each queue gets reserved slice of CPU time to schedule amongst its processes (e.g.: 80% for hi-priority, 20% low)

Policy #5: Multilevel Queues

An Example:



Policy #5: Multilevel Queues

In UNIX:

32 queues

→ 0...7: **Assignment Project Exam Help**
System (kernel) queues
(e.g.: 0 reserved for swapper)
<https://tutorcs.com>

→ 8...31: **WeChat: cstutorcs**
User-level queues
Processes move from queue to queue
1. **Nice**: Downward
2. Aging: Upward

Policy #5: Multilevel Queues

Variation: “Multilevel Feedback Queue”

- *Process can move between queues*

MFQ Scheduler Must Have Following Parameters:

1. *Number of queues*
2. *Scheduling algorithms for each queue*
3. *Scheduling algorithm between queues (or CPU time-slices)*
4. *Upgrade process (e.g.: Aging)*
5. *Downgrade process*
6. *Initial queue determination process*

<https://tutorcs.com>

WeChat: cstutorcs

Scheduling Design 101

Suppose Scheduling Goal Is: **Fairness**

(i.e.: Want to turnaround time to be proportional to burst time)

Accomplished Using MFQ's instead of Burst Prediction

A: *Queues*

<https://tutorcs.com>

“Reverse
aging”

1. *High priority: RR with low quantum*
2. *Low priority: RR with high quantum*

Process gets one go on each queue, getting lower priority as it gets older

Policy #5: Multilevel Queues

An Example of MFQ:

Q_0^{75} : RR, quantum = 75; *initial queue of entering process*

Q_1^{150} : RR, quantum = 150; *process ages to Q_1 after Q_0*

Q_2 : FCFS

Fixed priority

<https://tutorcs.com>

Interqueue Scheduling: FCFS (Serve Q_i when $Q_0 \dots Q_{i-1}$ empty)

[WeChat: cstutorcs](#)

<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
P ₁	125	0
P ₂	50	0
P ₃	500	0
P ₄	175	0
P ₅	200	250
P ₆	50	250

Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75$;	P_1	125	0
		P_2	50	0
Q_1^{150} :	$RR, \text{ quantum} = 150$;	P_3	500	0
		P_4	175	0
Q_2 :	$FCFS$	P_5	200	250
		P_6	50	250

Assignment Project Exam Help

<https://tutorcs.com>


Time = 0:

WeChat: cstutorcs

Q_0^{75} : $P_1^{125} \rightarrow P_2^{50} \rightarrow P_3^{500} \rightarrow P_4^{175}$

Q_1^{150} : —

Q_2 : —



Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75$;	P_1	50	0
		P_2	0	0
Q_1^{150} :	$RR, \text{ quantum} = 150$;	P_3	425	0
Q_2 :	$FCFS$	P_4	100	0
		P_5	200	250
		P_6	50	250

Assignment Project Exam Help

<https://tutorcs.com>

Time = 0:

WeChat: cstutorcs

Q_0^{75} : $P_1^{125} \rightarrow P_2^{50} \rightarrow P_3^{500} \rightarrow P_4^{175}$

Q_1^{150} : —

Q_2 : —



Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75$;	P_1	50	0
		P_2	0	0
Q_1^{150} :	$RR, \text{ quantum} = 150$;	P_3	425	0
		P_4	100	0
Q_2 :	$FCFS$	P_5	200	250
		P_6	50	250

Assignment Project Exam Help

<https://tutorcs.com>


Time = 0:

WeChat: cstutorcs

Q_0^{75} :

Q_1^{150} : $P_1^{50} \rightarrow P_3^{425} \rightarrow P_4^{100}$

Q_2 : —



Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75;$	P_1	50	0
		P_2	0	0
Q_1^{150} :	$RR, \text{ quantum} = 150;$	P_3	425	0
Q_2 :	$FCFS$	P_4	100	0
		P_5	200	250
		P_6	50	250

Assignment Project Exam Help

<https://tutorcs.com>

Time = 275:

WeChat: cstutorcs

Q_0^{75} : $P_5^{200} \rightarrow P_6^{50}$

Q_1^{150} : $P_1^{50} \rightarrow P_3^{425} \rightarrow P_4^{100}$

Q_2 : —

Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75$;	P_1	50	0
		P_2	0	0
Q_1^{150} :	$RR, \text{ quantum} = 150$;	P_3	425	0
		P_4	100	0
Q_2 :	$FCFS$	P_5	200	250
		P_6	50	250

Assignment Project Exam Help

<https://tutorcs.com>


Time = 275:

WeChat: cstutorcs

Q_0^{75} : $P_5^{200} \rightarrow P_6^{50}$

Q_1^{150} : $P_1^{50} \rightarrow P_3^{425} \rightarrow P_4^{100}$

Q_2 : —



Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75$;	P_1	50	0
		P_2	0	0
Q_1^{150} :	$RR, \text{ quantum} = 150$;	P_3	425	0
		P_4	100	0
Q_2 :	$FCFS$	P_5	200	250
		P_6	50	250

Assignment Project Exam Help

<https://tutorcs.com>


Time = 275:

WeChat: cstutorcs

Q_0^{75} : $P_5^{200} \rightarrow P_6^{50}$

Q_1^{150} : $P_1^{50} \rightarrow P_3^{425} \rightarrow P_4^{100}$

Q_2 : —



Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75$;	P_1	50	0
		P_2	0	0
Q_1^{150} :	$RR, \text{ quantum} = 150$;	P_3	425	0
		P_4	100	0
Q_2 :	$FCFS$	P_5	125	250
		P_6	0	250

Assignment Project Exam Help

<https://tutorcs.com>


Time = 275:

WeChat: cstutorcs

Q_0^{75} : $P_5^{200} \rightarrow P_6^{50}$

Q_1^{150} : $P_1^{50} \rightarrow P_3^{425} \rightarrow P_4^{100}$

Q_2 : —



Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75$;	P_1	50	0
		P_2	0	0
Q_1^{150} :	$RR, \text{ quantum} = 150$;	P_3	425	0
		P_4	100	0
Q_2 :	$FCFS$	P_5	125	250
		P_6	0	250

Assignment Project Exam Help

<https://tutorcs.com>


Time = 400:

WeChat: cstutorcs

Q_0^{75} : —

Q_1^{150} : $P_1^{50} \rightarrow P_3^{425} \rightarrow P_4^{100} \rightarrow P_5^{125}$

Q_2 : —



Policy #5: Multilevel Queues


		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, \text{ quantum} = 75$;	P_1	50	0
		P_2	0	0
Q_1^{150} :	$RR, \text{ quantum} = 150$;	P_3	425	0
Q_2 :	$FCFS$	P_4	100	0
		P_5	125	250
		P_6	0	250

Time = 400:

Q_0^{75} : —

Q_1^{150} : $P_1^{50} \rightarrow P_3^{425} \rightarrow P_4^{100} \rightarrow P_5^{125}$

Q_2 : —




Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	<i>RR, quantum = 75;</i>	P ₁	0	0
		P ₂	0	0
Q_1^{150} :	<i>RR, quantum = 150;</i>	P ₃	275	0
Q_2 :	<i>FCFS</i>	P ₄	0	0
		P ₅	0	250
		P ₆	0	250

Time = 400:

Q_0^{75} : —
 Q_1^{150} : P₁⁵⁰ → P₃⁴²⁵ → P₄¹⁰⁰ → P₅¹²⁵
 Q_2 : —



Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	<i>RR, quantum = 75;</i>	P ₁	0	0
		P ₂	0	0
Q_1^{150} :	<i>RR, quantum = 150;</i>	P ₃	275	0
Q_2 :	<i>FCFS</i>	P ₄	0	0
		P ₅	0	250
		P ₆	0	250

Assignment Project Exam Help

<https://tutorcs.com>


WeChat: cstutorcs

Time = 825:

Q_0^{75} : —

Q_1^{150} : —


Q_2 : P₃²⁷⁵



Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	$RR, quantum = 75;$	P ₁	0	0
		P ₂	0	0
Q_1^{150} :	$RR, quantum = 150;$	P ₃	275	0
Q_2 :	$FCFS$	P ₄	0	0
		P ₅	0	250
		P ₆	0	250

Time = 825:

 $Q_0^{75}:$ — $Q_1^{150}:$ —
$$Q_2: \quad P_3^{275}$$


Policy #5: Multilevel Queues

		<u>Process</u>	<u>Burst time</u>	<u>Arrival time</u>
Q_0^{75} :	<i>RR, quantum = 75;</i>	P ₁	0	0
		P ₂	0	0
Q_1^{150} :	<i>RR, quantum = 150;</i>	P ₃	0	0
		P ₄	0	0
Q_2 :	<i>FCFS</i>	P ₅	0	250
		P ₆	0	250

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Time = 1100:

Q_0^{75} : —

Q_1^{150} : —

Q_2 : —

Policy #5: Multilevel Queues

Process	Arrival time	Completion time	Turnaround Time
P ₁ (125)	0	450	450
P ₂ (50)	0	125	125
P ₃ (500)	0	1100	1100
P ₄ (175)	0	700	700
P ₅ (200)	250	825	575
P ₆ (50)	250	400	150

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Observe:

1. *Average turnaround time: $3100 / 6 = 516.7$*
2. *Shorter jobs tend to have shorter turnaround times*

CPU Scheduler Summary

- *First Example of Resource Management (Sharing)*

Scheduling Metrics

- *CPU utilization, throughput, turnaround time, wait time, response time*
- *also care about: fairness, dispatch time*

Scheduling Policies

<i>Policy</i>	<i>Type</i>	<i>+</i>	<i>-</i>
<i>FCFS</i>	<i>N</i>	<i>• easy implementation</i>	<i>• convoy effect</i>

CPU Scheduler Summary

- *First Example of Resource Management (Sharing)*

Scheduling Metrics

- *CPU utilization, throughput, turnaround time, wait time, response time*
- *also care about: fairness, dispatch time*

Scheduling Policies

<i>Policy</i>	<i>Type</i>	<i>+</i>	<i>-</i>
<i>FCFS</i>	<i>N</i>	<i>• easy implementation</i>	<i>• convoy effect</i>
<i>SJF</i>	<i>N,P</i>	<i>• waiting, turnaround time</i>	<i>• starvation</i> <i>• need to predict CPU bursts</i>

CPU Scheduler Summary

- *First Example of Resource Management (Sharing)*

Scheduling Metrics

- *CPU utilization, throughput, turnaround time, wait time, response time*
- *also care about: fairness, dispatch time*

Scheduling Policies

<i>Policy</i>	<i>Type</i>	<i>+</i>	<i>-</i>
<i>FCFS</i>	<i>N</i>	<i>• easy implementation</i>	<i>• convoy effect</i>
<i>SJF</i>	<i>N,P</i>	<i>• waiting, turnaround time</i>	<i>• starvation</i> <i>• need to predict CPU bursts</i>
<i>Priority</i>	<i>N,P</i>	<i>• importance of process considered</i>	<i>• waiting, turnaround time</i> <i>• starvation</i>

CPU Scheduler Summary

- *First Example of Resource Management (Sharing)*

Scheduling Metrics

- *CPU utilization, throughput, turnaround time, wait time, response time*
- *also care about: fairness, dispatch time*

Scheduling Policies

<i>Policy</i>	<i>Type</i>	<i>+</i>	<i>-</i>
<i>FCFS</i>	<i>N</i>	<i>• easy implementation</i>	<i>• convoy effect</i>
<i>SJF</i>	<i>N,P</i>	<i>• waiting, turnaround time</i>	<i>• starvation</i> <i>• need to predict CPU bursts</i>
<i>Priority</i>	<i>N,P</i>	<i>• importance of process considered</i>	<i>• waiting, turnaround time</i> <i>• starvation</i>
<i>Round-Robin</i>	<i>P</i>	<i>• fairness</i> <i>• response time</i>	<i>• turnaround time</i> <i>• context switch overhead</i>

CPU Scheduler Summary

- *First Example of Resource Management (Sharing)*

Scheduling Metrics

- *CPU utilization, throughput, turnaround time, wait time, response time*
- *also care about: fairness, dispatch time*

Scheduling Policies

<i>Policy</i>	<i>Type</i>	<i>+</i>	<i>-</i>
<i>FCFS</i>	<i>N</i>	<i>• easy implementation</i>	<i>• convoy effect</i>
<i>SJF</i>	<i>N,P</i>	<i>• waiting, turnaround time</i>	<i>• starvation</i> <i>• need to predict CPU bursts</i>
<i>Priority</i>	<i>N,P</i>	<i>• importance of process considered</i>	<i>• waiting, turnaround time</i> <i>• starvation</i>
<i>Round-Robin</i>	<i>P</i>	<i>• fairness</i> <i>• response time</i>	<i>• turnaround time</i> <i>• context switch overhead</i>
<i>MLQ</i>	<i>P</i>	<i>• can mix and match priority with other policies</i>	<i>• complex implementation</i> <i>• inherits weaknesses of scheduling policies it uses</i>